

PHOTONIC COMPUTING SYSTEM

Technical Design Document

Silicon Photonics + TFLN Modulator Architecture

Version 1.0 | February 20, 2026 | CONFIDENTIAL

49,585 TOPS

19.71 Tbps

953 TOPS/W

150x

Peak Throughput

Aggregate BW

Efficiency

Speedup vs GPU

Table of Contents

§	Section	Page
1.	Introduction & System Overview	3
2.	Architecture Overview	3
3.	Photonic Core Module	4
4.	TFLN Modulator Subsystem	5
5.	PCIe Interface & Host Integration	6
6.	FPGA Hybrid Integration	7
7.	WDM Optical Fabric	7
8.	REST API Reference	8
9.	Performance Metrics	9
10.	PCB Stack-Up & Manufacturing	10
11.	FEA Simulation Results	10
12.	Test Plan Summary	11
13.	Security & Reliability	11
14.	Glossary	12

1. Introduction & System Overview

The **LightRails AI Photonic Computing System** is a ground-up silicon-photonics accelerator platform that replaces electronic matrix arithmetic with optical interference, achieving sub-nanosecond matrix-vector products at energy costs orders of magnitude below GPU equivalents. This document is the authoritative Technical Design Document (TDD) for the system, covering hardware architecture, software interfaces, performance data, manufacturing constraints, and test methodology.

Document Scope

- Photonic core: Clements MZI mesh with ring-resonator amplitude weighting and SCL in-situ learning.
 - TFLN modulator subsystem: MZM, ring modulators, EO switches at 400 G – 800 G data rates.
 - PCIe Gen 5 host interface: DMA engine, MMIO register map, multi-board cluster.
 - FPGA hybrid integration: logic + photonic co-scheduling.
 - WDM optical fabric: 64-channel C-band multiplexing.
 - REST API server (Flask), Gerber/CNC/KiCad file services, FEA solver.
-

2. Architecture Overview

The system follows a layered architecture separating the optical compute plane from the electronic control and I/O plane. A PCIe Gen 5 link bridges the host CPU and the photonic accelerator board; an FPGA provides programmable glue logic for scheduling, calibration, and low-latency control.

Layer	Component	Technology
Host Software	Flask REST API Python SDK	Python 3.9 / NumPy / Numba
Host Interface	PCIe Gen 5 x16 DMA Engine	8 ch. DMA, MMIO @ 0xF0000000
Control Plane	Hybrid FPGA-Photonic Scheduler	1 M logic cells, 512-wide mesh
Compute Plane	Photonic Matrix Multiplier (128x128)	Clements MZI + Ring Resonators
Modulation	TFLN MZM + Ring Modulators	X-cut LiNbO ₃ , V _π < 2 V
Optical Fabric	WDM Multiplexer (64 ch.)	C-band 1530–1565 nm, 100 Gbps/ch
PCB / Packaging	4-layer FR-4, HHHL PCIe card	Gerber / KiCad design files

3. Photonic Core Module

3.1 Photonic Matrix Multiplier

The core compute element implements a **Clements decomposition MZI mesh** that realises an arbitrary $N \times N$ unitary matrix in $O(N^2)$ MZI stages. Matrix-vector multiplication is performed fully in the optical domain with a propagation latency of ~10 ps (single light-transit).

Parameter	Value
Matrix dimension (N)	128×128
Number of MZI elements	$N(N-1)/2 = 8,128$
Phase encoding range	$0 - 2\pi$ per MZI (θ, ϕ)
Propagation latency	10 ps (optical)
Base throughput	~16,384 TOPS
Operating wavelength	1550 nm
Q-factor (resonators)	15,000 (default)

3.2 Ring-Resonator Amplitude Weighting

Each MZI stage is preceded by a micro-ring resonator that applies a Lorentzian amplitude weight, boosting extinction ratio and SNR. Weights are computed via a vectorised binomial moment expansion (order $M=6$) across the full resonator array — $O(M \cdot N)$ with zero Python loops.

3.3 Statistical Congruential Learning (SCL)

SCL is an in-situ adaptive calibration loop that tunes phase settings without back-propagation. It uses a Knuth/MMIX-style LCG ($a=6364136223846793005$) and the **pigeonhole optimisation**: resonators are partitioned into $B = \lceil \sqrt{N} \rceil$ bins, requiring only $O(\sqrt{N})$ random draws per step instead of $O(N)$, guaranteeing full coverage by the pigeonhole principle.

SCL Parameter	Value
LCG multiplier (a)	6,364,136,223,846,793,005
LCG increment (c)	1,442,695,040,888,963,407
Modulus (m)	2^{32}
SCL bins (B)	$\lceil \sqrt{\text{num_mzi}} \rceil \approx 91$ for 128×128
Default epochs	50
Default lr	0.02
Training samples	24 per stream session

4. TFLN Modulator Subsystem

Thin-Film Lithium Niobate (TFLN) modulators are used for all high-speed optical I/O. TFLN offers a 3x lower $V\pi$ compared to silicon EO modulators, enabling PAM4/PAM8 links at 400 G – 800 G with sub-watt drive power.

4.1 Material Properties

Property	Value
Ordinary refractive index (n_o)	2.211 @ 1550 nm
Extraordinary index (n_e)	2.138 @ 1550 nm
EO coefficient r_{33}	30.8 pm/V
EO coefficient r_{11}	8.6 pm/V
Propagation loss (TE)	0.27 dB/cm
Propagation loss (TM)	0.30 dB/cm
$\chi^{(2)}$ (SHG)	27 pm/V
Thermo-optic dn/dT	$3.0 \times 10^{-4} \text{ K}^{-1}$

4.2 Mach-Zehnder Modulator Specs

Parameter	Spec
Interaction length	5,000 μm (active), 15 mm (link)
Electrode gap	10 μm (photonic core), 6 μm (link)
Wafer cut	X-cut (r_{33} mode)
Bandwidth	100 GHz (3 dB)
Insertion loss	3.5 dB
Extinction ratio	25 dB
$V\pi$ (API spec)	2.5 V
Energy/bit @ 400G PAM4	< 2 pJ/bit

4.3 Photonic Link Performance

Link Config	Data Rate	Reach	Power	Energy/bit
400G PAM4 (8-lane)	400 Gbps	2.0 km	~0.8 W	2.0 pJ/bit
800G PAM4 (8-lane)	800 Gbps	2.0 km	~1.5 W	1.9 pJ/bit
800G PAM8 (short)	800 Gbps	0.5 km	<1.5 W	<2 pJ/bit

5. PCIe Interface & Host Integration

The photonic accelerator board connects to the host via a **PCIe Gen 5 x16** slot, providing ~492 GB/s of bidirectional bandwidth. An 8-channel DMA engine handles zero-copy transfers; memory-mapped I/O at base address **0xF000_0000** exposes all photonic component control registers.

Parameter	Value
PCIe generation	Gen 5 (32 GT/s per lane)
Lane width	x16
Raw bandwidth	512 Gbps
Effective BW (128b/130b)	~492 Gbps (~61.5 GB/s)
DMA channels	8 independent
MMIO base address	0xF000_0000
Form factor	HHHL (Half-Height Half-Length)
Power consumption	75 W (board only)
Optical ports	16 (4 laser sources)

5.1 MMIO Register Map

Register	Offset	Description
CONTROL	0x0000	System control (reset, enable)
STATUS	0x0004	System status flags
LASER_POWER	0x0008	12-bit DAC laser power
MODULATOR_BIAS	0x000C	Modulator bias voltage
PHASE_SHIFTER_0	0x0010	Phase shifter ch.0 (16-bit)
PHASE_SHIFTER_1	0x0014	Phase shifter ch.1 (16-bit)
WDM_CHANNEL_SELECT	0x0018	Active WDM channel select
DMA_CONTROL	0x0028	DMA engine control
PERFORMANCE_COUNTER	0x002C	TOPS measurement counter

6. FPGA Hybrid Integration

A hybrid FPGA-photonic co-processor handles workloads that benefit from reconfigurable digital logic alongside optical matrix operations. The FPGA provides 1 M logic cells and schedules partitioning between digital and photonic compute paths.

Parameter	Value
FPGA logic cells	1,000,000
FPGA compute	1,500 GFLOPS
Photonic mesh size	512-wide
Photonic throughput	250 TOPS
Optical I/O	3.2 Tbps
Combined throughput	~250 TOPS + 1.5 TFLOPS
Total power (hybrid)	25 W

7. WDM Optical Fabric

The wavelength-division multiplexing subsystem carries 64 independent data channels across the C-band (1530 – 1565 nm). Each channel operates at 100 Gbps, yielding 6.4 Tbps aggregate throughput. Ring resonators tune each channel for drop/add routing.

Parameter	Value
Number of channels	64
Wavelength range	1530 – 1565 nm (C-band)
Channel spacing	~0.56 nm (~70 GHz)
Per-channel rate	100 Gbps
Aggregate bandwidth	6.4 Tbps
Per-channel power	-3 dBm (nominal)
Resonator type	Micro-ring (FSR = 20 nm)

8. REST API Reference

The Flask application server exposes a comprehensive REST API at <http://0.0.0.0:5001>. All data endpoints return JSON. Simulation endpoints accept POST with a JSON body.

Endpoint	Method	Description
/api/performance	GET	Aggregate system performance metrics
/api/matrix_multiplier	GET	Matrix multiplier specs & throughput
/api/matrix_multiplier/resonator	GET	Ring-resonator weight statistics
/api/matrix_multiplier/scl_stream	GET	SSE stream: real-time SCL training epochs
/api/matrix_multiplier/scl_train	POST	Blocking SCL train (legacy)
/api/matrix_multiplier/scl_reset	POST	Reset MZI phases & resonances
/api/matrix_multiply	POST	Run matrix-vector multiply simulation
/api/wdm	GET	WDM system performance
/api/wdm/channels	GET	Per-channel wavelength & power
/api/fft	GET/POST	FFT specs or run simulation
/api/fft_transform	POST	Execute optical FFT simulation
/api/pcie	GET	PCIe board raw performance
/api/pcie_board	GET	PCIe board (frontend format)
/api/pcie_transfer	POST	Simulate DMA transfer
/api/hybrid	GET	FPGA-photonic hybrid specs
/api/cluster	GET	64-board cluster aggregate metrics
/api/tfln/modulator	GET	TFLN MZM specs
/api/tfln/link	GET	TFLN link specs (800G)
/api/tfln/link_400g	GET	400G link parameters
/api/tfln/link_800g	GET	800G link parameters
/api/tfln/comparison	GET	TFLN vs. Si-photonics comparison
/api/tfln/plots	GET	Generate TFLN characterisation plots
/api/gerber/files	GET	List Gerber PCB files
/api/gerber/view/<file>	GET	Parse & return single Gerber layer
/api/gerber/layers	GET	All Gerber layers for visualisation
/api/gerber/projections	GET	Orthographic PCB projection views
/api/cnc/files	GET	List CNC G-code files
/api/cnc/view/<file>	GET	Parse G-code commands

/api/3d_models	GET	List STL/OBJ 3D model files
/api/vlsi/layout	GET	VLSI die layout (multi-layer Gerber)
/api/fea/simulate	POST	Run optical mode FEA simulation
/api/kicad/visualize	GET	Parse KiCad PCB geometry
/api/kicad/fea	POST	Thermal + EM FEA on KiCad board
/api/execute_workload	POST	Partition & execute DNN workload

9. Performance Metrics

System-level performance figures are derived from the live API (/api/performance) combining matrix multiplier, WDM, FFT, PCIe, hybrid, and cluster subsystems.

Metric	Single Board	64-Board Cluster	Unit
Peak Throughput	~49,586	~500	TOPS
Aggregate BW	~19.71	12.8	Tbps
Energy Efficiency	953.6	~2,500	TOPS/W
Speedup vs. GPU	150x	150x	—
Matrix Mult. TOPS	16,384	1.05 M	TOPS
WDM Bandwidth	6.4	409.6	Tbps
PCIe Bandwidth	~492	~31,488	Gbps
FFT Throughput	50	3,200	TOPS
Total Power	52	200	W

10. PCB Stack-Up & Manufacturing

The LightRails AI photonic board uses a cost-optimised **4-layer FR-4** PCB targeting standard fabrication at JLCPCB / PCBWay. The original 12-layer Rogers design was rationalised to reduce BOM cost by ~60%.

Layer	Type	Net	Thickness
F.Cu (Top)	Signal	—	35 µm
In1.Cu	Plane	GND	35 µm
In2.Cu	Plane	+3V3	35 µm
B.Cu (Bot)	Signal	—	35 µm

Gerber file set (24 files): copper layers, solder mask, silkscreen, board outline, and NC drill files (via 0.3 mm min). An optimised Gerber ZIP is available at [gerber_optimized/](#).

11. FEA Simulation Results

The FEA solver computes optical mode profiles and thermal maps using a finite-difference method across a 2D cross-section of the photonic waveguide.

Simulation Type	Parameter	Simulated Result
-----------------	-----------	------------------

Optical mode	Waveguide: w=0.5 μm, h=0.22 μm, λ=1.55 μm	Single-mode TE■■
Thermal (PCB)	Max component temp (TFLN modulator)	45.2 °C
Thermal (PCB)	Average board temperature	32.1 °C
EM – Impedance	50 Ω trace target / actual	50.0 / 49.8 Ω
EM – Crosstalk	Differential pair isolation	-42.3 dB
EM – Return loss	Signal path return loss	-18.5 dB

12. Test Plan Summary

Production testing follows the *Production Test Plan Final* document. Key stages are summarised below.

Test Stage	Criteria	Tool
Unit – Photonic Core	$\text{multiply}()$ output norm $\neq 0$; throughput > 1 TOPS	pytest / NumPy
Unit – TFLN MZM	$V\pi \in [1.5, 5]$ V; BW > 50 GHz	pytest
Unit – WDM	64 channels, BW = 6.4 Tbps	pytest
Integration – API	All <code>/api/*</code> routes return HTTP 200	requests + pytest
Integration – SCL	Loss decreasing over 50 epochs	SSE stream parser
System – Throughput	Live TOPS $\geq 49,000$	<code>/api/performance</code>
System – PCB FEA	Max temp < 85 °C; EM crosstalk < -35 dB	<code>fea_integration.py</code>
Acceptance – BER	BER < 10^{-12} @ 400G PAM4, 2 km	BERT tester

13. Security & Reliability

- **API Security:** Flask debug mode is disabled in production; server binds to 0.0.0.0:5001 — firewall rules must restrict external access.
- **Input Validation:** All POST endpoints validate JSON fields and clamp numeric parameters before passing to simulation kernels.
- **Error Handling:** Uncaught exceptions in FEA/Gerber routes return HTTP 500 with a JSON error body; tracebacks are logged to `app.log`.
- **Optical Safety:** Laser power is DAC-limited to 100 mW max via the LASER_POWER register (12-bit, full-scale = 4095).
- **Thermal Protection:** FEA hotspot alerting flags components exceeding 60 °C in simulation; hardware OTP fuse at 85 °C.
- **MTBF Estimate:** Optical components (LDs, photodetectors) rated > 100,000 hours; PCIe re-link on error via PERST#.

14. Glossary

Term	Definition
BER	Bit Error Rate — probability of an incorrectly decoded bit.
Clements mesh	Optimal rectangular MZI decomposition for universal unitaries (Clements et al., 2016).
DMA	Direct Memory Access — zero-CPU data transfer between host RAM and device.
EO	Electro-Optic — optical effect driven by applied electric field.
FSR	Free Spectral Range — wavelength spacing between resonator modes.

LCG	Linear Congruential Generator — fast pseudo-random number generator.
MMIO	Memory-Mapped I/O — device registers mapped into CPU address space.
MZI	Mach-Zehnder Interferometer — beam splitter + phase shift + combiner.
PAM4	4-level Pulse Amplitude Modulation — 2 bits per symbol.
SCL	Statistical Congruential Learning — in-situ adaptive phase calibration algorithm.
TFLN	Thin-Film Lithium Niobate — EO platform with low $V\pi$ and high bandwidth.
TOPS	Tera Operations Per Second — unit of compute throughput.
WDM	Wavelength Division Multiplexing — simultaneous multi-channel optical transmission.
$V\pi$	Half-wave voltage — voltage required for π phase shift in a MZM.